



Supervised Learning: Introduction to Supervised Learning

AAA-Python Edition



Plan

- 1- Supervised Learning
- 2- Classification
- 3- Regression
- 4- Features
- 5- Linear Regression
- 6- Polynomial Regression



1- Supervised Learning

Machine Learning

Machine learning

Build

Mathematical Model

with

Tunable parameters

Fit

Observed Data: described
by a **set of** features

Predict

New Data



1- Supervised Learning

Supervised Learning

Supervised learning

Build

Mathematical Model

with

Tunable parameters

Fit

Labeled Observed Data

Fitting == **minimize** the **cost** function
(**difference** between **True** and **Predicted** labels)

Predict the Labels

Predicting == **Apply** the **obtained mathematical** model on the **New Data**



2- Classification

Unsupervised Learning

Unsupervised learning

Build

Mathematical Model

with

Tunable parameters

Fit

Not Labeled Observed Data

Predict the Labels

Predicting = **Extract** information about **New Data**



2- Classification

Classification

Supervised learning: Classification

Build

Mathematical Model

with

Tunable parameters

Fit

Labeled Observed Data: The Labels are **categories**

Predict

Predict the **known categories** for the New Data



2- Classification

Scikit-Learn

- Scikit-learn is a python library used for data mining and data analysis
- It is build on : "Numpy", "Scipy", and "Matplotlib" libraries.
- It will be used for both classification and regression by importing the corresponding modules.
- It will also be used for feature extraction using : "sklearn.feature_extraction" module
- The full documentation about scikit-learn, is available at its homepage: <http://scikit-learn.org/stable/index.html>



3- Regression

Classification Problems

Classifying images

Fit

Observed Data: **x-ray** images. Labels == **yes** or **no**.
Yes == Presence of Tumor, No == Absence of Tumor

Predict if there is a tumor or not

New Data: Not labeled x-ray images.

Classifying text

Observed Data: **words** in sentences. Labels == **verb**, **subject**, or **adjective**

Predict if the category of the word

New Data: Not labeled words



3- Regression

Regression

Supervised learning: Regression

Build

Mathematical Model

with

Tunable parameters

Fit

Labeled Observed Data: The Labels are **continuous** quantities

Predict

Predict the **New quantities** for the New Data



3- Regression

Regression Problems

House pricing

Fit

Observed Data: **houses** described by a set of **characteristics**.
Labels == **prices** of the houses

Predict houses price

New Data: a set of houses described by the same set of characteristics (with different values), but they don't have the price indicated.

Photometric redshift

Observed Data: **galaxies** described by their **brightness** at several **wavelengths**. Labels == **distances** of the galaxies

Predict the distance

New galaxies without the distance information



4- Features

Introduction

Machine Learning

Fit

Observed Data: the data must be described by a set of **characteristics: numerical values**

Predict if there is a tumor or not

The new Data: must be described by the same set but with different values

These characteristics, aren't always in a numerical format. They can be:

Categories

Text

Images

They have to be transformed in a numerical format



4- Features

Categorical features

Categories



One-hot-encoding

The Feature 3 is transformed into 3 other features each one representing the presence of a category belonging to feature 3

```
from pandas import DataFrame as DF, Series as S
```

```
# we assume that the features are contained in dataframe dataF  
# the feature: "Feat3" is a categorical feature (not numeric)
```

```
dataF = DF([{"Feat1": 545, "Feat2": 3, "Feat3": "Cat1"}, {"Feat1": 362, "Feat2": 2, "Feat3": "Cat3"}, {"Feat1": 1005, "Feat2": 5, "Feat3": "Cat2"}], index=["house1", "house2", "house3"])  
dataF
```

```
# The DictVectorizer from sklearn will be used for the one-hot-encoding  
from sklearn.feature_extraction import DictVectorizer  
# convert the dataframe to a dictionary values  
myDict = dataF.to_dict('records')  
# create an instance of a DictVectorizer  
vec = DictVectorizer(sparse=False, dtype=int)  
# transform the category data  
trDict = vec.fit_transform(myDict)  
# convert the dictionary to a dataframe (just for visualization purpose)  
DF.from_dict(trDict)
```

myDict

```
[{'Feat1': 545, 'Feat2': 3, 'Feat3': 'Cat1'},  
{ 'Feat1': 362, 'Feat2': 2, 'Feat3': 'Cat3'},  
{ 'Feat1': 1005, 'Feat2': 5, 'Feat3': 'Cat2'}]
```

	Feat1	Feat2	Feat3
house1	545	3	Cat1
house2	362	2	Cat3
house3	1005	5	Cat2

		0	1	2	3	4
house1	545	3	1	0	0	
house2	362	2	0	0	1	
house3	1005	5	0	1	0	

[By Amina Delali]



4- Features

Text features

Text

Word count

TF – IDF:

term frequency–inverse
document frequency

Word count

```
# a multilines string
text= '''this is a simple text
we will count and we will see this'''
# converting each line to a list element
textT = text.split("\n")
# import CountVectorizer for word count encoding
from sklearn.feature_extraction.text import CountVectorizer
# create an instance of CountVectorizer
vec = CountVectorizer()
# create the encoding
wc = vec.fit_transform(textT)
# create the corresponding dataframe
DF(wc.toarray(), columns=vec.get_feature_names())
```

Each line will represent
a row. And each word
will represent a feature.
Features values in each
row will be the count of
the corresponding word
in the corresponding
line

['this is a simple text', 'we will count and we will see this']

	and	count	is	see	simple	text	this	we	will
0	0	0	1	0	1	1	1	0	0
1	1	1	1	0	1	0	0	1	2

Appears 0 times in line 0
Appears 2 times in line 1



4- Features

Text features

Text

TF – IDF

Word count

TF – IDF:
term frequency–inverse
document frequency

```
# import TfidfVectorizer for frequency-inverse document frequency encoding
from sklearn.feature_extraction.text import TfidfVectorizer
# create an instance of CountVectorizer
vec = TfidfVectorizer(norm=None)
# create the encoding
wc = vec.fit_transform(textT)
# create the corresponding dataframe
DF(wc.toarray(), columns=vec.get_feature_names())
```

We didn't normalize the results

$$idf(t) = \ln\left(\frac{1+n_d}{1+df(d,t)}\right) + 1$$

$tf-idf(t,d) = tf(t,d) * idf(t)$ $n_d = \text{total number of documents}$

$tf(t,d) = \text{number of } \times \text{the term } t \text{ occurs} \in \text{document } t$

$df(d,t) = \text{number of documents that contain the term } t$

$$\begin{aligned} &= 1^* \\ &(\ln[(1+2)/(1+1)] + 1) \\ &= 1.405465 \end{aligned}$$

['this is a simple text', 'we will count and we will see this']

	and	count	is	see	simple	text	this	we	will
0	0.000000	0.000000	1.405465	0.000000	1.405465	1.405465	1.0	0.000000	0.000000
1	1.405465	1.405465	0.000000	1.405465	0.000000	0.000000	1.0	2.81093	2.81093

[By Ami]



4- Features

Images

Image

Use pixel values as
features values:
number of features ==
number of pixels


```
# importing digits dataset sample
from sklearn.datasets import load_digits
# loading digits dataset sample
digits = load_digits()
# the dataset has pixel values for 1797 images,
# each image is coded in 8*8 pixel values
# indicating grayscale values
print(digits.images.shape)
# for example, the 35th image represent a "6"
digits.images[34]
```

```
(1797, 8, 8)
array([[ 0.,  0.,  0.,  5., 14.,  2.,  0.,  0.],
       [ 0.,  0.,  1., 13., 11.,  0.,  0.,  0.],
       [ 0.,  0.,  5., 16.,  2.,  0.,  0.,  0.],
       [ 0.,  0.,  6., 15.,  5.,  0.,  0.,  0.],
       [ 0.,  1., 15., 16., 15., 11.,  1.,  0.],
       [ 0.,  2., 13., 14.,  1., 12.,  9.,  0.],
       [ 0.,  0.,  4., 16.,  7., 13.,  9.,  0.],
       [ 0.,  0.,  0.,  5., 16., 15.,  3.,  0.]])
```


```
1 # plotting the 35th image ( the digit 6)
2 %matplotlib inline
3
4 import matplotlib.pyplot as plt
5 plt.imshow(digits.images[34])
```

digits.data[34]

Is just a
reshape
of images



```
4 myList = (([10]*8)*2+ ([0]*3+[13]*2+[0]*3)* 6
5 myT= np.array(myList).reshape(8,8)
6 plt.imshow(myT)
```



[By Amina Delali]



5- Linear Regression

Simple Linear Regression (with 1 Feature)

Simple Linear Regression : 1 feature

Build

Mathematical Model== a
Line, modeled by:
 $Y = \mathbf{a} x + \mathbf{b}$

with

a: the slope
b: the intercept

Fit

Labeled Observed Data: described by 1 feature: **x**
The labels are **y** values. Fitting == find **a** and **b** that minimize the
difference between the real labels: **y** and the estimated ones.

Predict

Predict the **New y** for the New not labeled **x** values



5- Linear Regression

Example

```
1 # importing form linear_model module, LinearRegression
2 from sklearn.linear_model import LinearRegression
3 # create an instance of the model
4 model = LinearRegression()
5 # fitting the model using x and y values
6 model.fit(np.array(x).reshape(-1,1), np.array(y).reshape(-1,1))
7 # values of a (coef) and b(intercept) after fitting
8 print("a == the value in ", model.coef_)
9 print("b == the value in ", model.intercept_)
10 # new values of x : we create this values to generate the
11 # line representing the model
12 newX = np.linspace(1,20,1000)
13 # predict the labels for newX
14 newY = model.predict(np.array(newX).reshape(-1,1))
15 # plotting the old and the new values
16 plt.scatter(x, y)
17 plt.plot(newX, newY, color="g");
```

The model to train

Initialization of a linear regressor instance


Train the model

Parameters of the model

Predict y for new values: actually this new x values are used to draw the model

```
a == the value in [[3.029415]]
b == the value in [-1.97286737]
```

The line representing the model





6-Polynomial Regression

Linear regression basis Polynomial function (1 feature)

Linear regression basis Polynomial function (1 feature)

Build

Mathematical Model== a Linear curve, modeled by:

$$Y = a_1x + a_2x^2 + \dots + a_nx^n + b$$

with

The parameters: a_1, a_2, \dots, a_n, b

N: the degree of the polynomial model

Fit

Labeled Observed Data: described by 1 feature: x
From that x , new **polynomial features** are generated: x^2, x^3, \dots, x^n .

Predict

The new x values must be **transformed** first into polynomial feature, before applying the model.



6-Polynomial Regression

Example


```
1 from sklearn.preprocessing import PolynomialFeatures
2 from sklearn.pipeline import make_pipeline
3 xArr= np.array(x).reshape(-1,1)
4 y2Arr= np.array(y2).reshape(-1,1)
5 # the order of the polynomial features
6 order =11
```

The model is the same, we will just generate new features: **polynomial features**

```
polyObj= PolynomialFeatures(order,include_bias=False)
myNewFeatures=polyObj.fit_transform(xArr)
# instantiate a linear regressor model
myModel = LinearRegression()
# fitting the model
myModel.fit(myNewFeatures, y2Arr)
# the calculated parameters
for i in range(1,order+1):
    print("a"+ str(i)+" == "+ str(myModel.coef_[0,i-1]))
print ("b == ",myModel.intercept_)

# new values of x : to visualize the model
newX = np.linspace(1,20,1000)
newXArr= np.array(newX).reshape(-1,1)
# corresponding polynomial features for the new values
newPolyX =polyObj.fit_transform(newXArr)
# predict the labels for newX
newy2 = myModel.predict(newPolyX)
# plotting the old values and the visualization of the model
plt.scatter(x, y2)
plt.plot(newX, newy2, color="g");
```

```
a1 == -0.057443431666221394
a2 == -0.10775669970224205
a3 == -0.15702923169271918
a4 == -0.09393489084681118
a5 == 0.12136012879617761
a6 == -0.04165097865144431
a7 == 0.007255471629826568
a8 == -0.0007421930578640621
a9 == 4.655613420523336e-05
a10 == -1.7661168388830744e-06
a11 == 3.727072246228341e-08
a12 == -3.3641874008782935e-10
b == [0.81724983]
```





6-Polynomial Regression

Example with a pipeline


```
from sklearn.pipeline import make_pipeline
# create the pipeline
myPipeline= make_pipeline(polyObj,myModel)
myPipeline.fit(xArr, y2Arr)
# the calculated parameters
for i in range(1,order+1):
    print("a"+ str(i)+" = "+ str(myModel.coef_[0,i-1]))
print ("b = ",myModel.intercept_)

# predict the labels for newX
newy2 = myPipeline.predict(newXArr)
# plotting the old values and the model visualization
plt.scatter(x, y2)
plt.plot(newX, newy2, color="q");
```

This pipeline will:

- 1- generate polynomial features from the data.
- 2- apply the regression model to the new data

The newX values are generated in a way that the model can be plotted in a linear form : it describes the model





References

- Joshi Prateek. Artificial intelligence with Python. Packt Publishing, 2017.
- Jake VanderPlas. Python data science handbook: essential tools for working with data. O'Reilly Media, Inc, 2017.



Thank you!

FOR ALL YOUR TIME