



Supervised Learning: Multivariable Regressor & Classifiers

AAA-Python Edition



Plan

- 1- Preprocessing Data
- 2- A single variable regressor
- 3- A multivariable regressor
- 4- Regularization
- 5- Logistic Regression Classifier
- 6- Confusion matrix

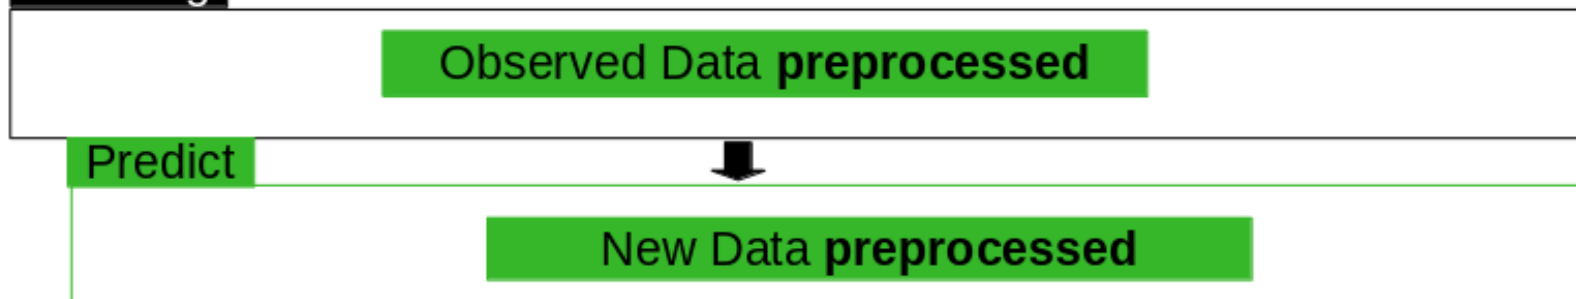


1- Preprocessing Data

Introduction

- Before going further, we have to focus in one important task: **preprocessing data**.
- Before starting the learning (and the predicting) process, the used **data must be prepared** or **transformed**.

Learning



- The different processes that can be applied to the data are:
 - Binarization
 - Mean removal
 - Scaling
 - Normalization

Binarization

Define a **threshold**Values $>$ **threshold** \implies 1Values \leq **threshold** \implies 0

myD

	Feature1
s1	-0.626009
s2	1.556353
s3	0.450549
s4	0.898856
s5	-0.191403

```

1 # the value of the threshold
2 th = 1
3 # instaciate a Binarizer object with that threshold
4 myBin = preprocessing.Binarizer(threshold=th)
5 # transform the feature's values with that Binarizer
6 myD_B =myBin.transform(myD.Feature1.values.reshape(-1,1))

```

```

array([[0.],
       [1.],
       [0.],
       [0.],
       [0.]])

```

Mean removal

Calculate the "mean" \implies Values - mean

```
myD_mr = preprocessing.scale(myD,with std=False)
```

Previous mean == **0.502482**New mean: *new mean \approx 0*

```

array([[ -0.11428357],
       [ -0.34343912],
       [  0.13638221],
       [  1.27876145],
       [ -0.95742096]])

```

Min Max Scaling

Features with different scales of values

Select a range

Apply the formula

Determine min and max for each feature

myD

	Feature1	Feature2	Feature3
s1	-0.626009	100	9
s2	1.556353	1000	7
s3	0.450590	500	-3
s4	0.898856	200	1
s5	-0.191403	80	55

```
1 # initiate a minmax scaler instance
2 myScaler=preprocessing.MinMaxScaler(feature_range=(0, 1))
3 # scale the features in myD
4 myD_s = myScaler.fit_transform(myD)
```

```
array([[0.        , 0.02173913, 0.20689655],
       [1.        , 1.        , 0.17241379],
       [0.49331825, 0.45652174, 0.        ],
       [0.6987223 , 0.13043478, 0.06896552],
       [0.19914478, 0.        , 1.        ]])
```

The transformation is done using this formula:

(from Google colab help)

$$X_std = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))$$

$$X_scaled = X_std * (max - min) + min$$

min, max = feature_range

Ex: $X[s1, Feature2]_{std} = (100 - 80) / (1000 - 80) = 0.02173913$
 $X[s1, Feature2]_{scaled} = 0.02173913 * (1 - 0) + 0 = 0.02173913$

Normalization

Least Absolute Deviations: L1

Sum of **absolute** features values for each sample row will be == 1

```
1 # Normalize myD using L1 normalization
2 myD_l1 = preprocessing.normalize(myD, norm='l1')
```

$$|-5.71040582e-03| + |9.12192288e-01| + |8.20973059e-02| == 1$$

```
array([[ -5.71040582e-03,  9.12192288e-01,  8.20973059e-02],
       [ 1.54314927e-03,  9.91516237e-01,  6.94061366e-03],
       [ 8.95003420e-04,  9.93146120e-01, -5.95887672e-03],
       [ 4.45201136e-03,  9.90595014e-01,  4.95297507e-03],
       [-1.41579269e-03,  5.91753604e-01,  4.06830603e-01]])
```

Least Squares: L2

Sum of the **square of** features values for each sample row will be == 1

```
1 # Normalize myD using L2 normalization
2 myD_l2 = preprocessing.normalize(myD, norm='l2')
```

$$(-6.23476844e-03)^2 + (9.95955081e-01)^2 + (8.96359573e-02)^2 == 1$$

```
array([[ -6.23476844e-03,  9.95955081e-01,  8.96359573e-02],
       [ 1.55631299e-03,  9.99974290e-01,  6.99982003e-03],
       [ 9.01163413e-04,  9.99981594e-01, -5.99988957e-03],
       [ 4.49417844e-03,  9.99977401e-01,  4.99988701e-03],
       [-1.97154737e-03,  8.24040323e-01,  5.66527722e-01]])
```



2- A Single Variable Regressor

Training and testing set

Learning : regression (with 1 feature)

build

Mathematical model

with

parameters

Fit : determine the value of parameters

Observed **Labeled Data X**: described by 1 feature and the labels **y**.

The data is split into : training and testing set

Parameters determined using training set

Prediction is made using the testing set

Compute the error using predicted and test labels

```
from sklearn.model_selection import train_test_split
# splitting the data into test and training sets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

```
# predict using testing set
y_test_pred = myModel.predict(x_test)
```

```
# create an instance of a Linear regressor
myModel = linear_model.LinearRegression()
# train the model using the training sets
myModel.fit(x_train, y_train)
```

```
# Compute mean absolute error
print("Mean absolute error =", sm.mean_absolute_error(y_test, y_test_pred))
# Compute mean squared error
print("Mean squared error =", sm.mean_squared_error(y_test, y_test_pred))
```

Mean of absolute differences between test labels and predicted labels

[By Amina Delali]

Mean of square of the differences between test labels and predicted labels



2- A Single Variable Regressor

Loading and saving models

```
1 # creating a new file
2 f= open("myModel.pkl", 'wb')
3 # saving the model to that file
4 pickle.dump(myModel, f)
```

Saving the model


```
1 # opening (loading) the file
2 f= open("myModel.pkl", 'rb')
3 #loading the model from that file
4 myLModel = pickle.load(f)
```

Loading the model

```
# try other splits
x1,x2,y1,y2= train_test_split(x, y, test_size=0.3)
```

Real
(original)
values

Predicted
(estimated)
values



30% of the values will be used for test (x_2 , y_2) and 70% should be used for training (x_1, y_1)

Since the model is Already trained, we will just use the data to test the model



3- A Multivariable Regressor

Multiple features

Learning : regression (with multiple features)

build

Mathematical model

with

parameters

Fit : determine the value of parameters

Observed **Labeled Data X:**

- described by **1 feature**, then **new features** are **generated**: like **polynomial features**
- described by **multiple features**: they are **used as they are**
- described by **multiple features**, and **new features** are **generated**

The data here is described by
3 features

```
# move to AAA-Ped-Week4 folder
%cd AAA-Ped-Week4
# loading the data:
# "col0""col1""col2" are x features values: 3 features
# and "col3" are the labels y
myD3 = pd.read_csv("A3P-w4-data_multivar_regr.txt",header=-1,prefix="col")
```

	col0	col1	col2	col3
0	2.06	3.48	7.21	15.69
1	6.37	3.01	7.27	15.34
2	1.18	1.20	5.42	0.66
3	7.37	3.81	-1.95	38.37
4	6.16	1.39	7.39	9.96

[By Amina Delali]



3- A Multivariable Regressor

Simple Linear Regression (with **multiple** features)

Learning : Simple linear regression (with multiple features)

build

$$\text{Mathematical model}== \\ y = a_1 * x_1 + a_2 * x_2 + ... + a_n * x_n + b$$

with **n+1** parameters: a_1 and b

Fit : determine the value of a_1 and b

Observed **Labeled** Data: **X** described by **multiple** features
the **labels**: **y**

```
1 # Features values
2 myX= myD3.iloc[:,3]
3 # labels
4 myY = myD3.iloc[:,3]
5 myY = myY.values.reshape(-1,1)
6 # train again the linear regression model
7 myModel.fit(myX,myY)
8 print("ai are in ",myModel.coef_)
9 print ("b is in ",myModel.intercept_)
```

Selecting the **3** first **columns**
corresponding to **3** features

The **fourth** column is the **labels** column

a_1, a_2 , and a_3

b is in [5.77068196]

ai are in [[1.97288257 3.91615283 -1.75291766]]

b



3- A Multivariable Regressor

Polynomial Basis Function (with **multiple** features)

Learning: Linear regression (with 3 features, order == 2)

build

$$\text{Mathematical model}==$$
$$y = a_1 * x_1 + a_2 * x_2 + a_3 * x_3 + a_4 * x_1^2 + a_5 * x_1 * x_2$$
$$+ a_6 * x_1 * x_3 + a_7 * x_2^2 + a_8 * x_2 * x_3 + a_9 * x_3^2 + b$$

with

10 parameters: a_i and b

Fit : determine the value of a_i and b

Observed **Labeled Data**: X described by **multiple features** +
Polynomial features are generated
the labels: y

```
from sklearn.preprocessing import PolynomialFeatures
# polynomial features generator of degree 2
polynomial = PolynomialFeatures(degree=2, include_bias=False)
# generate the new features
x_poly = polynomial.fit_transform(myX)

# fit the model
myModel.fit(x_poly, myY)
```

```
ai are in [[ 2.24816066e+00  4.01398972e+00 -1.95742341e+00 -3.11506706e-02
-3.60516775e-02  3.01335037e-02 -5.82113466e-03  2.13183791e-02
-4.29337992e-04]]
b is in [5.67788315]
```

[By Amina Delali]



4- Regularization

Introduction

Mathematical model **overfit** data == **Fits** very well to the training set. But, **doesn't perform** well on the **testing** or the **new data sets**

Introduce penalties to the model


Lasso Regression (L1 regularization): Penalizes the sum of absolute values (1-norms) of regression parameters: \mathbf{a}_1 and \mathbf{b}

Lasso Regression (L1 regularization): Penalizes the sum of squares (2-norms) of the regressions parameters: \mathbf{a}_1 and \mathbf{b}

Linear model without regularization

Ridge model

Lasso model





4- Regularization

Ridge Regression (L2 regularization)

Penalty formula


$$p = \alpha * \sum_{n=1}^N \theta_n^2$$

```
# a Ridge regressor: linear regressor with L2 regularization
regressor2 = Ridge(alpha=0.1)

# a Ridge regressor with polynomial features
myRModel= make_pipeline(PolynomialFeatures(17), regressor2)

2 myRModel.fit(x_tr, y_tr)
0 y_pred2 = myRModel.predict(x_plot)
```

```
1 y_test_pred = myRModel.predict(x_plot)
2 plt.scatter(x_te, y_te, color="red")
3 plt.plot(x_plot, y_pred2, color="orange")
4 plt.show()
```




[By Amina Delali]

The first linear model without regularization and the test values

Smoother and more fitting to the test data

```
y_test_pred = myModel2.predict(x_plot)
plt.scatter(x_te, y_te, color="red")
plt.plot(x_plot, y_pred, color="green")
#x_tes= np.linspace(0, 10, 1000)
plt.show()
```





4- Regularization

Lasso Regression (L1 regularization)


Penalty formula

$$p = \alpha * \sum_{n=1}^N |\theta_n|$$

```
# a Lasso regressor: linear regressor with L1 regularization  
regressor3 = Lasso(alpha=0.1)
```

```
# a Lasso regressor with polynomial features  
myLModel = make_pipeline(PolynomialFeatures(17), regressor3)  
myLModel.fit(x_tr, y_tr)  
y_pred3 = myLModel.predict(x_plot)
```

```
1 y_test_pred = myLModel.predict(x_plot)  
2 plt.scatter(x_te, y_te, color="red")  
3 plt.plot(x_plot, y_pred3, color="black")  
4 plt.show()
```




[By Amina Delali]

The first linear model without regularization and the test values

More smoother but less fitting to the test data

```
y_test_pred = myModel2.predict(x_plot)  
plt.scatter(x_te, y_te, color="red")  
plt.plot(x_plot, y_pred, color="green")  
#x_tes= np.linspace(0, 10, 1000)  
plt.show()
```





5- Logistic Regressor Classifier

Definition

Classification with a logistic regression

build

Mathematical model == a **logistic Regression defined by a regression model and the sigmoid function S:** with

$$s(z) = \frac{1}{1 + e^{-z}}$$

where z is :the prediction value obtained by a regression model

Tunable parameters:
The regression model
parameter

Fit

Observed **Labeled** Data: a defined cost function is used for the logistic Regression. The S values are the probabilities of the samples data belonging to a Certain class.

Predict the knows Categories

New Data: each sample is classed into one of the known categories. If the probability value is higher than a certain threshold, the sample Is affected to that class. And if it is not, it is affected to the other class.



5- Logistic Regressor Classifier

Classification with 2 classes


```
# instance of a logistic regressor
myModel_reg = linear_model.LogisticRegression()
# train the model
myModel_reg.fit(x2_train, y2_train)
# predict for test values:
y2_test_pred= myModel_reg.predict(x2_test)
```

The “+” marks **colors** represent the **predicted** classes. The “.” marks **colors** represent the **real** values

This sample was predicted to be in class2 (2) , but in fact, it belongs to class1 (1)

y axis :
second
feature
values

x axis : first feature values



```
1 # the coefficients calculated of the regression model
2 a = myModel_reg.coef_[0]
```

```
# the corresponding probabilities predicted the 3 first samples
y2_test_pr = myModel_reg.predict_proba(x2_test)[0:3]
y2_test_pr
```

```
array([[0.0145635  0.9854365 ],
       [0.99701667 0.00298333],
       [0.00170004 0.99829996]])
```

Probabilities for
belonging to class2

```
# the intercept of the regression model
b= myModel_reg.intercept_
```

Calculated
parameters
for the
regression
model

```
array([-0.12507469])
```

```
array([ 1.34836231, -2.05135853])
```




5- Logistic Regressor Classifier

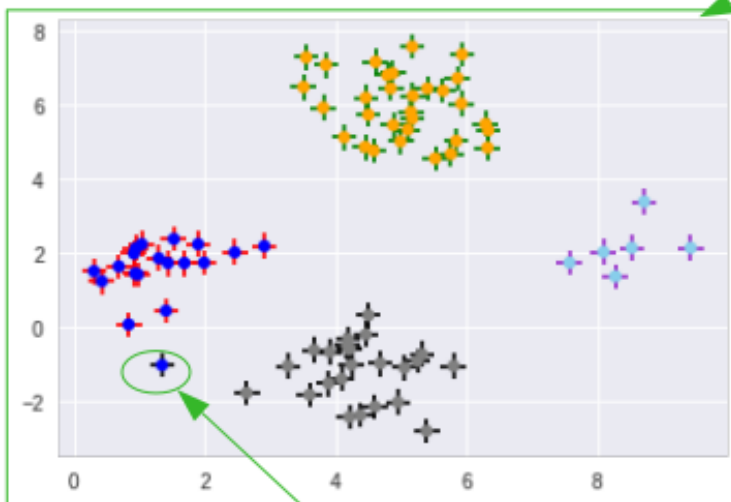
Classification with **more than 2 classes: One vs all**

- The logistic regression is run for each class against all the other classes (considered as one class).
- Select the class for which the probability of belonging is the **greatest**.

An other way for visualization (the code is from :Joshi Prateek. Artificial intelligence with Python. Packt Publishing, 2017 [link](#))


```
9 # instance of a logistic regressor
10 myModel_reg = linear_model.LogisticRegression()
11 # train the model
12 myModel_reg.fit(x2_train, y2_train)
13 # predict for test values:
14 y2_test_pred= myModel_reg.predict(x2_test)
```

Same code as logistic regression with 2 classes



[By Amina Delali]

Wrong predicted class



Predicted class

Real class



6- Confusion Matrix

Definition

Well
Predicted

True Negatives: number of samples predicted to be in class 0 and they are in class 0	False Positives: number of samples predicted to be in class 0 but they are in class 1
False Negatives: number of samples predicted to be in class 1 but they are in class 0	True Positives: number of samples predicted to be in class 1 and they are in class 1

Predicted
wrong

```
5 true_classes = [1,1,1,0,1,0,0,1,0,1,0]
7 pred_classes = [1,0,1,0,1,0,1,1,1,1,1]
8 # Create confusion matrix
9 myConfusionMat = confusion_matrix(true_classes, pred_classes)
```

2 samples are known to be in class 0, and they were predicted to be in class 0 : **True** (the good prediction) **Negatives** (prediction in class 0)

array([[2, 3],
 [1, 5]])

3 samples are known to be in class 0, and they were predicted to be in class 1 : **False** (the wrong prediction) **Positives** (prediction in class 1)



6- Confusion Matrix

More than 2 classes

In scikit-learn

Number of classes == $n+1$

Samples are known to be in (the rows number):

(0 , 0)	(0 , 1)	(0 , 2)	...	(0 , n)
(1 , 0)	(1 , 1)	(1 , 2)	...	(1 , n)
(2 , 0)	(1 , 1)	(2 , 2)	...	(2 , n)
...
(n , 0)	(n , 1)	(n , 2)	...	(n , n)

Samples predicted to be in (columns number)



6- Confusion Matrix

Confusion Matrix Report

Number of samples that truly belongs to class 4

Confusion matrix

```
array([[3, 2, 1, 1, 1],  
       [2, 6, 0, 0, 0],  
       [1, 2, 1, 0, 0],  
       [1, 0, 0, 1, 0],  
       [0, 1, 0, 0, 2]])
```

```
1 print(classification_report(true_classes, pred_classes))
```

	precision	recall	f1-score	support
0	0.43	0.38	0.40	8
1	0.55	0.75	0.63	8
2	0.50	0.25	0.33	4
3	0.50	0.50	0.50	2
4	0.67	0.67	0.67	3
avg / total	0.51	0.52	0.50	25

$0.43 == 3 / (3 + 2 + 1 + 1 + 0) ==$
True positives / (true positives / false positives)

$0.25 == 2 / (1 + 2 + 1 + 0 + 0) ==$
True Positives / (True Positives + False Negatives)

$0.63 == 2 * (0.55 * 0.75) / (0.55 + 0.75) == 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$



References

- Joshi Prateek. Artificial intelligence with Python. Packt Publishing, 2017.
- Jake VanderPlas. Python data science handbook: essential tools for working with data. O'Reilly Media, Inc, 2017.
- Scikit-learn.org. scikit-learn, machine learning in python. On-line at scikit-learn.org/stable/. Accessed on 03-11-2018



Thank you!

FOR ALL YOUR TIME