



# Unsupervised Learning: Dimensionality reduction

AAA-Python Edition



# Plan

- 1- Dimensionality reduction
- 2- Some Math
- 3- PCA
- 4- PCA in scikit-learn
- 5- Manifold Learning
- 6- Manifold Examples



# 1- Dimensionality reduction

## Objective

- **Dimensionality reduction** in **machine learning** is reducing the **number of features** of the training dataset.
- This reduction is necessary to:
  - Eliminate the **noise** from the data
  - Visualize the data in **2** or **3** dimensions
  - Speed up the learning process
  - Enhance the learning results by eliminating correlated features.
  - Eliminate unnecessary features.
  - Compress the data size.
- Two main approaches to dimensionality reduction are:
  - **Projection** : project the data into a lower dimensional space.
  - **Manifold**: suppose that the data in the higher dimension is just a manifold of a representation of the data in the lower dimension.



## 1- Dimensionality reduction

### Projection

- Sometimes the degree of the variation of the data is different from one dimension to another. So, for some features, the values can be very diverse, and for others, they can barely change.
- So we project the data into a lower dimension in order to keep only the most influential **information** ==> we define a mapping between the original data from the higher dimension to new data in a lower dimension.
- The most used technique to define this mapping, is **PCA** (**P**roincipal **C**omponent **A**nalysis) and its variations:
  - **Incremental** PCA
  - **Randomized** PCA
  - **Kernel** PCA



## Manifold

- Like we said earlier, we make the hypothesis that our data is created from a **manifold** of a data in a **lower dimension**. So, reducing it to this low dimension is like **straightening up** this manifold (or **unrolling it**).
- The different techniques used, are:
  - **MDS**: Multidimensional Scaling. Tries to preserve the distances between instances.
  - **LLE**: Locally Linear Embedding. Tries to preserve the relationship between a sample and its closets points.
  - **Isomap**: the samples will represent nodes of a graph. These nodes are connected to their closets neighbors. The algorithm tries to preserve the number of nodes in the shortest path connecting two nodes.



### Singular value decomposition

- It is the the decomposition of a matrix  $\mathbf{M}_{(m,n)}$  into **3** matrices:  $\mathbf{U}_{(m,m)}$ ,  $\mathbf{S}_{(m,n)}$ , and  $\mathbf{V}_{(n,n)}$ . Considering only **real** values, we have the following characteristics:
  - $\mathbf{M} = \mathbf{U} . \mathbf{S} . \mathbf{V}^T$  (  $\mathbf{V}^T$  is the transpose matrix of  $\mathbf{V}$  : value at  $i,j$  becomes at  $j,i$  )
  - $\mathbf{U} . \mathbf{U}^T = \mathbf{U}^T . \mathbf{U} = \mathbf{I}_{(m,m)}$  (the identity matrix)
  - $\mathbf{V} . \mathbf{V}^T = \mathbf{V}^T . \mathbf{V} = \mathbf{I}_{(n,n)}$
  - The diagonal (values with the same row and column indices) of  $\mathbf{S}$  are the **Singular values** of  $\mathbf{M}$ 
    - ➔ Singular values are the square roots of **eigenvalues**
    - ➔ The other values of  $\mathbf{S}$  are **zeros**.
  - The columns of  $\mathbf{U}$  are the **eigenvectors** of  $\mathbf{M} . \mathbf{M}^T$ .
  - The columns of  $\mathbf{V}$  are the **eigenvectors** of  $\mathbf{M}^T . \mathbf{M}$ .



### Eigenvectors, Eigenvalues

- Given  $\mathbf{A}_{(n,n)}$  a square matrix:
  - If  $\mathbf{A} \cdot \mathbf{V}_{(n)} = \lambda \cdot \mathbf{V}_{(n)}$  then:  $\mathbf{V}$  is an **eigenvector** and  $\lambda$  is its corresponding **eigenvalue**.
  - The above equation can be rewritten as follow:  $(\mathbf{A} - \lambda \mathbf{I}) \cdot \mathbf{V} = 0$
  - Several  $\lambda$  can solve the equation. For each  $\lambda$  lambda value, an eigenvector is computed.
- Example:
  - If  $\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$ 
    - Its eigenvalues will be: 1 , 3
    - And their corresponding eigenvectors will be:  $\begin{bmatrix} 1 \\ -1 \end{bmatrix}$  and  $\begin{bmatrix} 1 \\ 1 \end{bmatrix}$



### Standard Deviation

- The standard deviation  $\sigma$  measures how data is spread (or distant from the mean) . It is the **square root** of the **variance**.

- The variance is computed as follow:

$$\text{variance} = \frac{\sum_{i=1}^N (x_i - \mu)^2}{N}$$

- And the standard deviation:

$$\sigma = \sqrt{\text{variance}}$$

- To project data on new axis, we select the axis that preserve the maximum possible variance of the data. This way, most of the information is preserved.





## Definition

- It is a linear dimensionality reduction technique that project data using orthogonal axes (components) that preserve the maximum variance possible. One of the method used is singular value decomposition of the mean centered training data.
- As stated before the decomposition leads to **3** matrices. The vectors of the matrix  $V^T$  will be used to project the data. They are the “principal components”.
- Each component will conserve a certain amount of variance. The variance obtained after projection is the accumulation of the variances obtained by each component
- To project, we select a sufficient number of component to preserve the maximum of variance, then we apply the transformation (the projection), using only this number of vectors.
- The number of vectors will determine the dimension of the projection.



## 3- PCA in scikit-learn

### Example

```
# import necessary libraries
import numpy as np
from sklearn.datasets import load_iris
import sklearn.preprocessing as preprocess
```

```
1 # loading the data ( 4 feautres ==> 4 dimensions)
2 myIris = load_iris()
3 X= myIris.data
4 y=myIris.target
5
```

```
from numpy import mean
# centring the data
X_centred = preprocess.scale(X,with_std=False, axis=0)
X_centred = X - X.mean(axis=0)
# extracting the principal component
U, s, V = np.linalg.svd(X_centred)
# extracting the 3 first principal components
c1 = V.T[:,0]# V.T is the transpose of V
c2 = V.T[:,1]
c3 = V.T[:,2]
```

```
1 # compute the projection in a 3D dimension
2 C3 = V.T[:, :3]
3 X3d = X_centred.dot(C3)
```

Center the data to the mean, before applying the decomposition

The decomposition



To project, we multiply the centered data by the first selected component==> we will have a 3 dimensions projection



### 3- PCA in scikit-learn

## Results

3D projection



2D projection

Since our data was originally labeled (we don't use those label for decomposition), we used them for coloring the data.

And what is obvious, is that the data is clustered according to its classes.

Which proofs:

- that the clustering can in certain cases classify data.
- the decomposition preserved the most important amount of information.



## 4- Processing Data



### With matplotlib

```
1 # we are using matplotlib version 2.1.2
2 # it will be removed in version 3.1
3 from matplotlib.mlab import PCA as PCA2
4 my2PCA = PCA2(X, standardize=False)
5 results = my2PCA.project(X, minfrac=0.02)
6 fig = plt.figure()
7 plt.scatter(results[:,0], results[:,1], c=y, cmap = plt.cm.Set1)
```

It tells to only center the data, and to not standardize

It will drop all the axis with variance ratio  $< \text{minfrac}$ . In this case, it will only keep 2 axis.

Same results as in our previous implementation



[By Amina Delali]



## 4- Processing Data



### With sklearn

```
1 from sklearn.decomposition import PCA
2
3 fig = plt.figure()
4 pca = PCA(n_components = 2)
5 pca.fit(X)
6 X2d_2= pca.transform(X)
7
8
9 plt.scatter(X2d_2[:,0],X2d_2[:,1],c=y, cmap = plt.cm.Set1)
```

We have to select the number of components before transforming the data

As in matplotlib, we don't have to center the data

The reason of this inversion is that **sklearn** **flip** the **eigenvector's sign** before the projection : it apply the method **svd\_flip** on the vectors U and V in the fitting methods



Comparing with matplotlib we see that the directions are inverted



## 4- Processing Data

### Explained variance ratio

- The **correct number of components** can be defined by the **explained variance ratio** of each component.
- It is computed by the value of **explained variance** divided by the **sum of all variances**.
- The ratio of each component are **summed up** until a certain percentage is obtained.
- The **variances** can be computed from the **square** of the **singular values** in **S**

```
# the explained variance ratio (our implementation, 2D)
explained_ratio_2FirstC = (np.square(s[0]) + np.square(s[1])) / np.sum(np.square(s))
```

0.977685206318795

```
# explained variance ratio (with matplotlib, 2D)
EVR2 = my2PCA.fracs[0] + my2PCA.fracs[1]
```

```
1 # explained variance ratio (with sklearn, 2D)
2 EVR2 = np.sum(pca.explained_variance_ratio_)
```



## 5- Manifold Learning: LLE

### Algorithm

- **LLE** for **L**ocally **L**inear **E**mbedding. The algorithm consist of **3** major steps:
- **Step 1 - identifying the neighbors for each sample  $x_i$  from the data  $X_{(N,D)}$  (for N samples and D features) :**
  - Compute the distances of the other samples from  $x_i$
  - Select the **k** smallest distances.
- **Step 2 - for each sample  $x_i$  compute its neighbors weights:**
  - Create the matrix  $Z_{(k,D)}$  with the k samples rows from  $X_{(N,D)}$  corresponding to the neighbors of  $x_i$
  - Subtract  $x_i$  values from each row of  $Z_{(k,D)}$
  - Compute  $C_{(k,k)} = Z_{(k,D)} \cdot Z_{(D,k)}^T$  (  
in the original page it is inverted because of X and Z are transposed)
  - Compute the row **i** of the matrix  $W_{(N,N)}$  with:
    - ➔ Compute the weights in the one column vector  $w_{(k,1)}$  that solve the equation  $C_{(k,k)} \cdot w_{(k,1)} = 1_{(k,1)}$  (1 is a column vector with only 1 as values)



## Algorithm (Suite)

- For the samples  $\mathbf{j}$  that do not belong to each  $x_i$  neighbors, set the weights to  $\mathbf{0}$ .
- For each neighbor  $\mathbf{b}$  of  $x_i$  set the weight to:  $\mathbf{w}(\mathbf{p}) / \sum(\mathbf{w}_{(k,1)})$ . Where  $\mathbf{p}$  is the indices in  $\mathbf{w}$  corresponding to the  $\mathbf{b}$  neighbor of  $x_i$ .
- **Step 3 - reduce the dimensionality to  $d < D$  in a new matrix  $\mathbf{Y}_{(N,d)}$ :**
  - Compute the matrix  $\mathbf{M}_{(N,N)} = (\mathbf{I}_{(N,N)} - \mathbf{W}_{(N,N)})^T \cdot (\mathbf{I}_{(N,N)} - \mathbf{W}_{(N,N)})$
  - Select the  $d+1$  eigenvectors of  $\mathbf{M}_{(N,N)}$  corresponding to the  $d+1$  smallest eigenvalues. Order these eigenvectors according to the corresponding eigenvalues sorted in a **decreasing order**.
  - For each **column  $\mathbf{q}$**  in  $\mathbf{Y}$  set the values equal to the values of the  $q+1$  smallest eigenvector counting from the bottom (to discard the last eigenvector corresponding to the eigenvalue 0)







## 5- Manifold Learning: LLE

### Example

```
from sklearn.datasets import make_swiss_roll  
#from sklearn.datasets.samples_generator import make_swiss_roll  
X_swiss, color = make_swiss_roll(n_samples=1500)
```

N == 1500, D == 3



```
from sklearn import manifold  
X_r, err = manifold.locally_linear_embedding(X_swiss, n_neighbors=12,  
                                             n_components= 2)
```

[By Amina Delali]

LLE : k == 12, d == 2



### Algorithm

- There are two types of Multidimensional Scaling: classical (or metric) that tries to reproduce the original distances. The second one is non-metric (**NMDS**) that tries to reproduce only the rank of the distances.
- We will describe the algorithm of the classical method using the euclidean distance:
  - Compute the distances between all points, and form a matrix of those distances in a matrix **D**.
  - Compute the matrix **A** as follow:  $A(i,j) = -1/2 * D(i,j)^2$
  - Compute the matrix **B** as follow:  $B(i,j) = A(i,j) - A(i,.) - A(.,j) + A(.,.)$   
where:  $A(i,.)$  is the average of all  $A(i,j)$  for a selected  $i$   
 $A(.,j)$  is the average of all  $A(.,j)$  for a selected  $j$   
 $A(.,.)$  is the average of all values of  $A$
  - Find the **p** (the **new dimension, lesser** than the original dimension ) largest eigenvalues of  $B$ :  
and their corresponding normalized eigenvectors  $L_1, L_2, \dots, L_p$   
so that  $L_i^T \cdot L_i = \lambda_i$



### Algorithm (suite)

- Form the matrix **L** as follow:  $L = (L_1, L_2, \dots, L_p)$ . The new values (coordinates) are the **rows** of L.
- This method minimizes the value of the **Stress**
- The **stress** is a measure that can be used to find the optimal lower dimension. It is computed as follow:
  - $$\text{stress} = \sqrt{\frac{\sum_{i < j} (D(i, j) - \Delta(i, j))^2}{\sum_{i < j} D(i, j)^2}}$$
  - where:  $\Delta$  is the matrix of the distances of the new matrix L
- A stress with a value  $< 0.05$  is acceptable, below 0.01 is considered to be good.

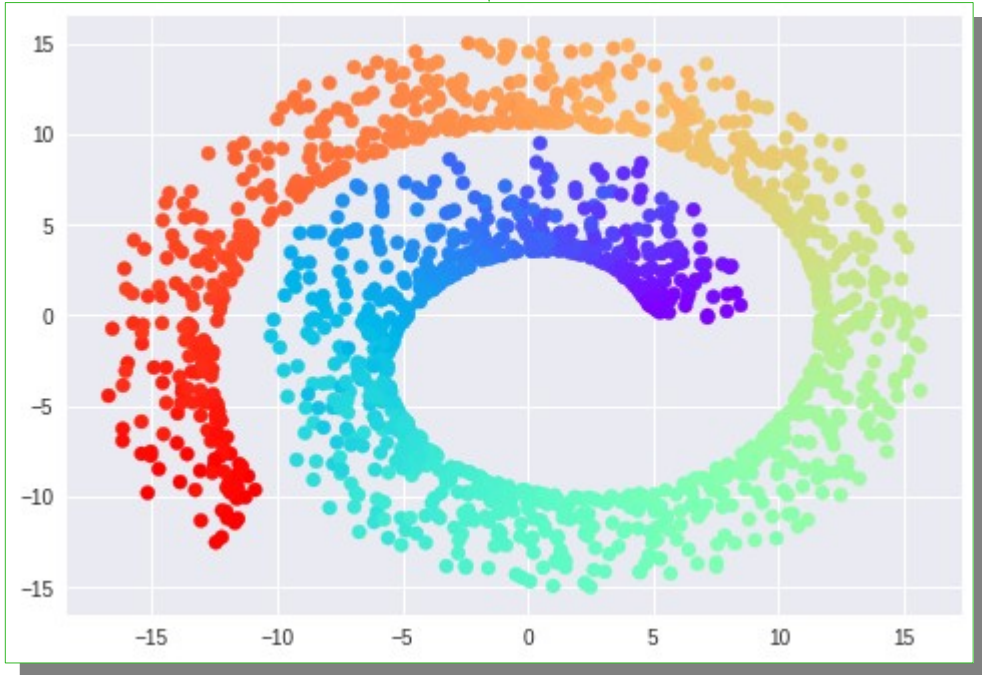


## 6-Polynomial Regression

### Example in Scikit-learn

```
1 from sklearn.manifold import MDS
2 embedding = MDS(n_components=2)
3 X_transformed_mds = embedding.fit_transform(X_swiss)
```

The results are completely different from the previous manifold technique. We see here, the goal is to keep the same original distances values as much as possible.





# References

- Aurélien Géron. Hands-on machine learning with Scikit-Learn and Tensor-Flow: concepts, tools, and techniques to build intelligent systems. O'Reilly Media, Inc, 2017.
- J. D. Hunter. Matplotlib: A 2d graphics environment. Computing In Science & Engineering, 9(3):90–95, 2007.
- NCSS Statistical Software. Multidimensional Scaling, ncss, llc edition.
- Scikit-learn.org. scikit-learn, machine learning in python. On-line at <https://scikit-learn.org/stable/>. Accessed on 03-11-2018.
- Jake VanderPlas. Python data science handbook: essential tools for working with data. O'Reilly Media, Inc, 2017.
- web.mit.edu. Singular value decomposition (svd) tutorial. On-line at [https://web.mit.edu/be.400/www/SVD/Singular\\_Value\\_Decomposition.htm](https://web.mit.edu/be.400/www/SVD/Singular_Value_Decomposition.htm). Accessed on 28-12-2018.
- wikipedia.org. Wikipedia, the free encyclopedia. On-line at <https://www.wikipedia.org/>. Accessed on 25-12-2018.



# Thank you!

FOR ALL YOUR TIME