



# **Artificial Neural Network: Introduction to ANN**

**AAA-Python Edition**



# Plan

- 1- Introduction to ANN
- 2- Perceptron
- 3- Single Layer Perceptron
- 4- Single Layer Perceptron examples
- 5- Multi Layer Perceptron
- 6- ANN topologies



## 1- Introduction to ANN

### Concept

- The idea of an **Artificial Neural Network** (ANN) is to build a model based on the way the human brain **learns new** things.
- It can be used in any type of **machine learning**. It learns by extracting the different underlying patterns in a given data.
- This extraction is performed by stages, called **layers**. Each layer, composed by a set of **neurons**, will identify a certain pattern. The following layer, will identify another **more complex** pattern, from its previous layer.
- The most common architecture is as follow: a first layer, has the training data as input. It is called the **input layer**. In the last one, the output of the neurons are the final output. It is called the **output** layer. The layers in between, are called **hidden layers**.
- From now, the term **neural network** will mean **artificial neural network**.



# 1- Introduction to ANN

## Neurolab

- **Neurolab** is Neural Network library for Python. It supports several types of neural networks.
- For installation: 

```
1 # installation of neurolab
2 !pip install neurolab
```
- Like other machine learning techniques, a neural network need to be **trained**. Can be **tested**. And will be used to **predict** results.
- Here is an example of how to use neurolab to create a neural network, and how to perform the fore-mentioned tasks:

```
1 import numpy as np
2 import neurolab as nl
3 # Create data
4 myInput = np.random.uniform(-0.5, 0.5, (10, 2))
5 # the labels correspond to the sum of two values
6 myLabels = (myInput[:, 0] + myInput[:, 1]).reshape(10, 1)
7 # Create network with 2 inputs, 5 neurons in hidden layer and 1 in output layer
8 myNN = nl.net.newff([[-0.5, 0.5], [-0.5, 0.5]], [5, 1])
9 # Train process
10 myErr = myNN.train(myInput, myLabels, show=15)
```

The samples are uniformly distributed from -0.5 (included) to 0.5 (excluded)

Length of the outer list  
Equals to the number of  
Neurons of the input layer

Input layer with **2** neurons,  
hidden layer with **5**  
neurons, and output layer  
with **1** neuron



# 1- Introduction to ANN

## Neurolab example details

- Data: **10** samples described by **2** features. The labels are the **sum** of the **2** features. In fact, the NN tries to model the **sum** function for values ranging from **-0.5** to **0.5**
- After creating the data, the steps were:
  - Create an instance of a neural network with specified number of layers and neurons (**nl.net.newff**)
  - Train the neural network (**myNN.train**)
  - Predict the output for the value [0.2,0.1] (**myNN.sim**)
  - Compute the test error (the true label is known: 0.2+0.1)

```
# Test and prediction process
pred= myNN.sim([[0.2, 0.1]])
# the result should be 0.3
testErr= np.abs(0.3-pred)
print ("Prediction=",pred)
print ("Test error = ",testErr)
```

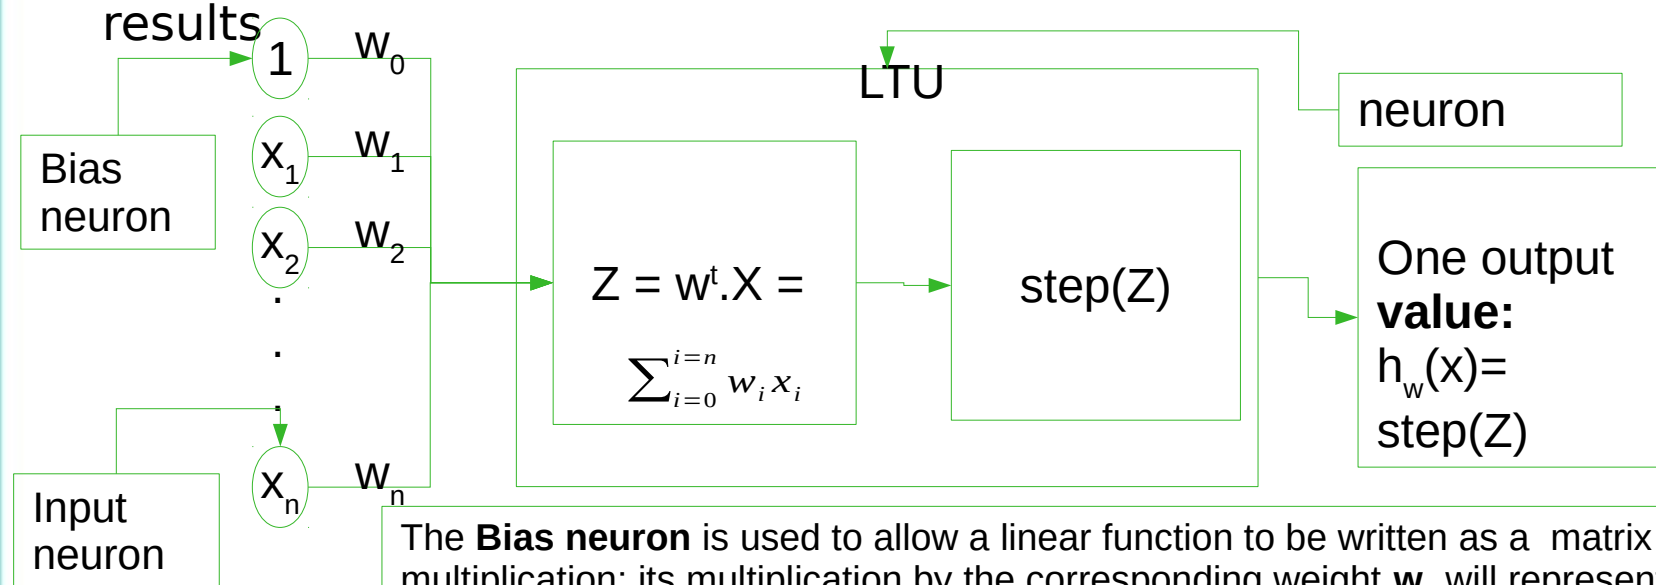
```
The goal of learning is reached
Prediction= [[0.29308242]]
Test error = [[0.00691758]]
```



## 2- Perceptron

### Definition

- The term **Perceptron** refers to an input layer of data features values, with forward weighted connections to an output **layer** of **one single neuron** , or of **multiple neurons**.
- One of the simplest form of a neuron is an **LTU**.
- **LTU**, for **L**inear **T**hreshold **U**nit, it is a component (neuron) that:
  - Computes a **weighted sum** of its inputs: a linear function
  - Applies a **step** function to the resulting sum, and **outputs** the results



The **Bias neuron** is used to allow a linear function to be written as a matrix multiplication: its multiplication by the corresponding weight  $w_0$  will represent the **intercept value  $b$** :  $(a_1x_1 + a_2x_2 + \dots + a_nx_n + b)$ .



### The functions

- The weighted sum function, is also called the **Propagation** function.
- The step function, can be :
  - A **non-linear** function, in this case, it will be called the **threshold activation function** (this is the case in an **LTU**). For example:
    - ➔ Heaviside step function:
$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$
    - ➔ Sign function:
$$\text{sgn}(z) = \begin{cases} -1 & \text{if } z < 0 \\ 0 & \text{if } z = 0 \\ 1 & \text{if } z > 0 \end{cases}$$
  - A **linear** function, simply called **activation function**. For example:
    - ➔ The **identity function**: which means that the value computed by the **propagation** function, **is** the **output** value of the neuron.
    - ➔ A **semi-linear** function, that is **monotonous** and **differentiable**. Also called **activation function**.



## 2- Perceptron

### Single Layer Perceptron

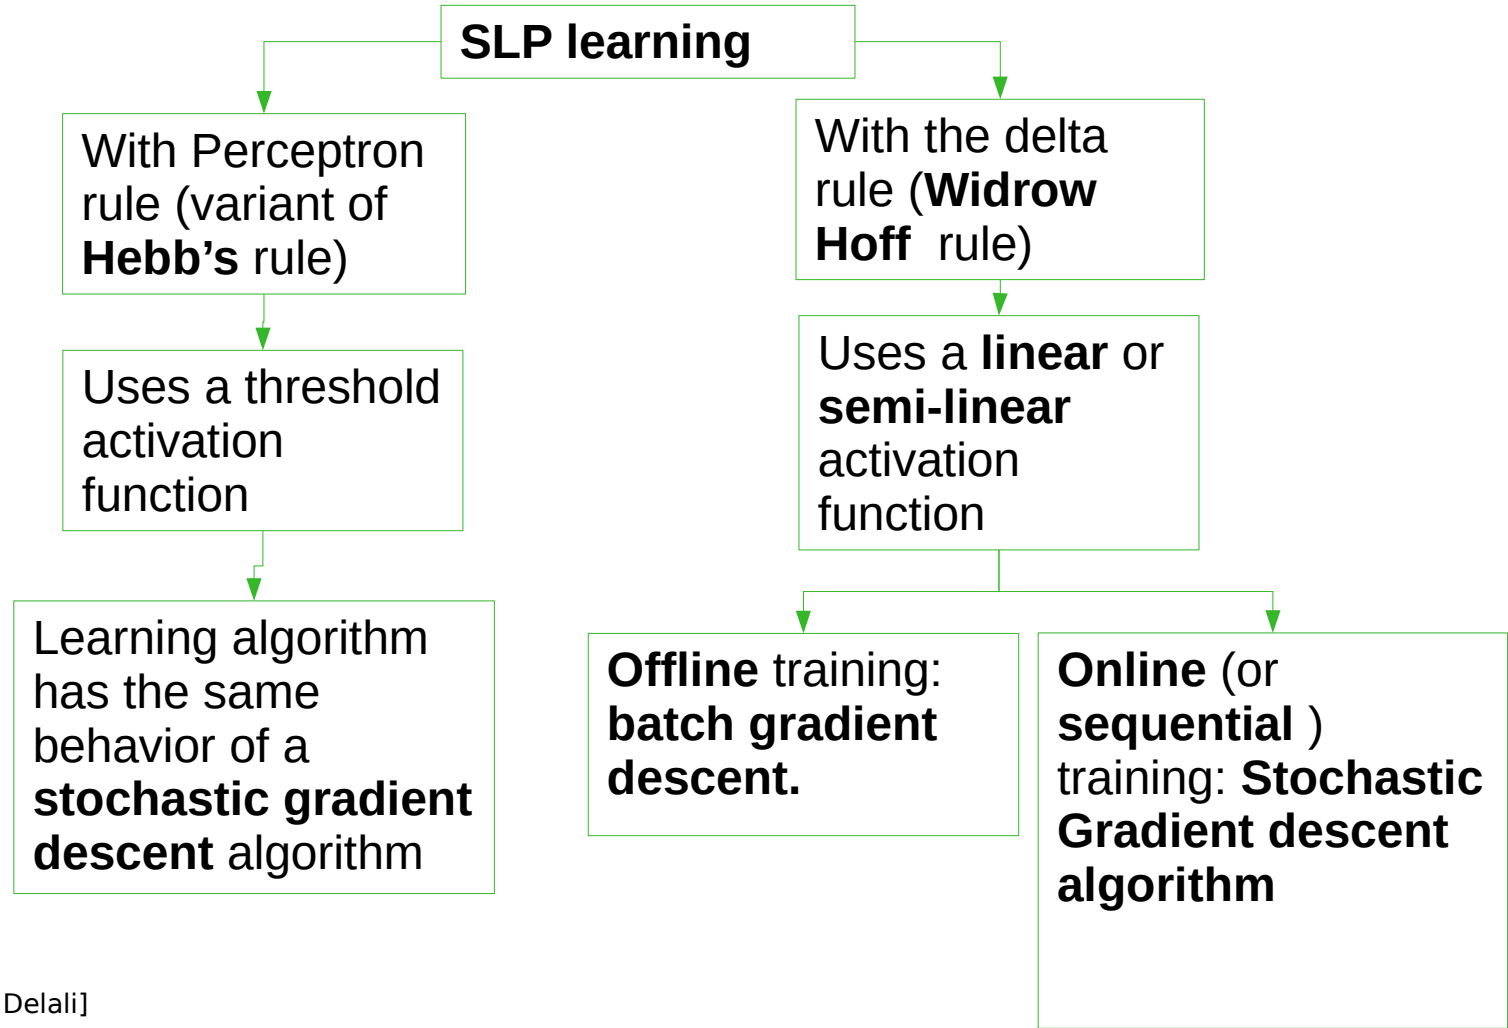
- A **Single Layer Perceptron (SLP)**, is simply a Perceptron, with only one layer (**without** counting the **input layer** ) .
- So, it is composed of an input layer and an output layer. The later one can have one ore more outputs. So, it can be used for binary and for multi-output classification
- Considering ANN in general, a **Perceptron** is considered as a **feedforward** neural network. We are going to talk about it in the next section.
- An **SLP** an apply **2** different kind for **rules** to learn: **the perceptron** rule, or the **delta rule**. Each of the rules is associated with a certain type of activation function. To apply the delta rule, we need the activation function to be differentiable.





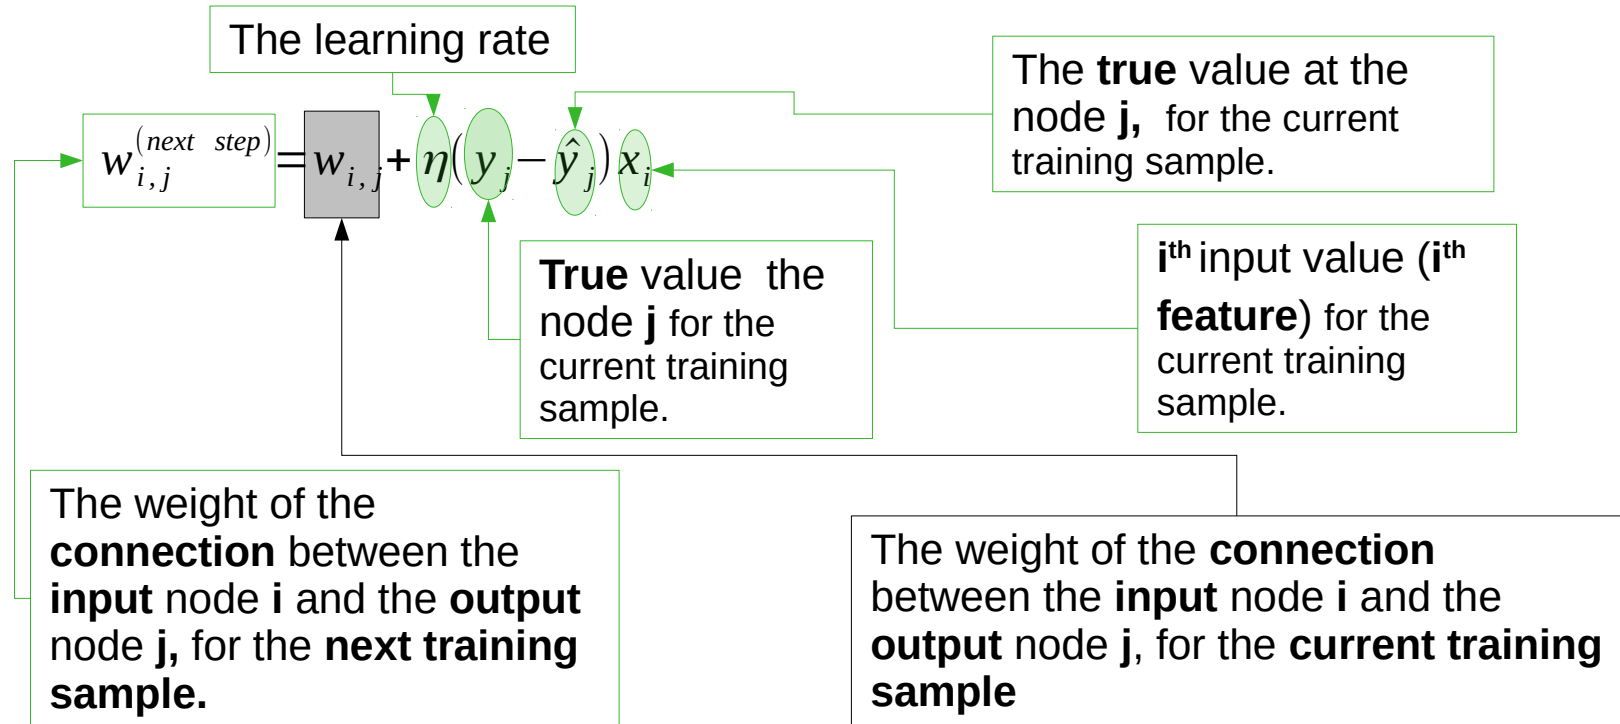
### 3- Single Layer Perceptron

## SLP Learning



## SLP Learning: Perceptron Rule

- To update the weights, the following formula is used:



- The concept is that each **wrong** prediction **reinforces** the **weight** corresponding to the **feature** that would **contributed** to the **correct prediction**. The computation of the weights is repeated until the samples are classified correctly.





### 3- Single Layer Perceptron

## Gradient Descent

- The concept of the **gradient descent** is:
  - With **initial parameters** of a model, predict an **output value**
  - Compute the gradient of the “**error**” (“loss”) function (function of the **parameters** of the learning model) at a certain point= the slope of the surface of that function at that point calculated by its **derivative** at that point.
  - Update the parameters in order to find the **local minima** by a step **proportional** to **the negative** of that **gradient**. (opposite direction ==> toward the local minima of the function). In the case of :
    - **A stochastic gradient descent**: with **one sample**, **predict** → **update** the parameters for the **next sample** → predict with the next sample with the new parameters
    - **A Batch gradient descent**: predict for all samples ==**1 epoch** → **update** the parameters → predict again with the new parameters
  - Repeat the process in order to **minimize** the error.



### 3- Single Layer Perceptron

- With the activation rule being **linear** or **semi-linear but differentiable**, the gradient descent is used to update the weights.
- The weights are updated as follow:  $w_{i,j}^{next} = w_{i,j} + \Delta w_{i,j}$ 
  - In general:  $\Delta w = \frac{-\eta \cdot \partial E}{\partial w}$

- In a case of a **linear activation function**, and a **Sum-Squared error function**

- **In Online training:** for a given sample:

$$\Delta w_{i,j} = \eta \cdot x_i \cdot (y_j - \hat{y}_j) = \eta \cdot \delta_j$$

This is why it is called the **delta rule**

The difference here with the perceptron rule is the type of activation function used

- **In Offline training:**

$$\Delta w_{i,j} = \eta \cdot \sum_{s \in X} x_i^s \cdot (y_j^s - \hat{y}_j^s) = \eta \cdot \sum_{s \in X} x_i^s \cdot \delta_j^s$$

The **input feature i** value of the sample **s**

The **true** label that should have the output neuron **j**, for the sample **s**

Learning rate

**Predicted** label for output neuron **j**, for the sample **s**



## 4- Single Layer Perceptron examples

### With sklearn: perceptron rule

```
#SLP using Perceptron class
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron
#load iris datasets
X, y = load_iris(return_X_y=True)
# instance of a perceptron: SLP with perceptron rule
# learning rate =0.1, maximum of epoch = 5000, without shuffle after
# each iteration
mySLP = Perceptron(eta0=0.1, max_iter=5000, tol=1e-3, shuffle=False)
mySLP.fit(X, y)
print ("The score of the classification (without shuffle) = ", mySLP.score(X, y) )

# instance of an other perceptron: SLP with perceptron rule
# learning rate =0.1, maximum of epoch = 20, with shuffle after
# each iteration
mySLPShuf = Perceptron(eta0=0.1, max_iter=20, tol=1e-3, shuffle=True)
mySLPShuf.fit(X, y)
print ("The score of the classification (with shuffle) = ", mySLPShuf.score(X, y) )
```

The fact that after each weight update the samples were **shuffled**, only **20** iterations were sufficient to reach convergence.

```
The score of the classification (without shuffle) =  0.6666666666666666
The score of the classification (with shuffle) =  0.82
```

Sklearn uses **Hinge loss function** with threshold = **0** to compute the error of the prediction

[By Amina Delali]

- Eta0==**0.1** → Learning rate = **0.1**
- max\_iter== **5000** → number of iterations if there is no convergence
- Shuffle == **False** → there is no shuffle after a weight is updated for the next sample



## 4- Single Layer Perceptron examples

### With sklearn: perceptron rule

```
from sklearn.linear_model import SGDClassifier
#SLP using SGDClassifier class
mySGDCl = SGDClassifier(loss="perceptron", learning_rate="constant", eta0=0.1, penalty = "none",
                        max_iter=20, shuffle=True, tol=1e-3, verbose=2, random_state=0)

mySGDCl.fit(X, y)

print ("The score of the classification (with shuffle) = ", mySGDCl.score(X, y) )
print("The number of classes = ", len(mySGDCl.classes_))
```

Use of stochastic  
gradient descent  
classifier class

Since there is **3**  
classes, and  
since the  
classification  
strategy used is  
**one vs all**, the  
classifier will run  
on **3 times**. Each  
time with **max  
epochs == 20**

→ Epoch 1

Norm: 0.83, NNZs: 4, Bias: 0.100000, T: 150, Avg. loss: 0.065987  
Total training time: 0.00 seconds.

-- Epoch 2

Norm: 0.83, NNZs: 4, Bias: 0.100000, T: 300, Avg. loss: 0.000000

→ Epoch 1

Norm: 1.27, NNZs: 4, Bias: 0.000000, T: 150, Avg. loss: 1.399133  
Total training time: 0.00 seconds.

-- Epoch 2

→ Epoch 1

Norm: 2.02, NNZs: 4, Bias: -0.400000, T: 150, Avg. loss: 1.023480  
Total training time: 0.00 seconds.

The score of the classification (with shuffle) = 0.82

The number of classes = 3



## 4- Single Layer Perceptron examples

### Example with neurolab: delta rule

```
f1_mi= X[:,0].min()
f1_ma= X[:,0].max()
f2_mi= X[:,1].min()
f2_ma= X[:,1].max()
f3_mi=X[:,2].min()
f3_ma= X[:,2].max()
f4_mi=X[:,3].min()
f4_ma= X[:,3].max()
# instantiate an SLP with: 4 input neurons for th 4 features
# with 3 output neurons : to represent the 3 classes
# The learning rule is the Delta rule, the activation function is the
# SoftMax function
mySLPDelta = nl.net.newp([[f1_mi,f1_ma],[f2_mi,f2_ma],[f3_mi,f3_ma],[f4_mi,f4_ma]],3,
                           transf=nl.trans.SoftMax())
```

The number of samples= 150  
The number of Features = 4

Since we have **4** features → we need an **input layer** with **4 neurons**. And since we have **3 classes**, and since we will use **SoftMax** activation function (ideal for multi-class classification), we will need **3** output neuron:  
Class 0 coded as 1 0 0  
Class 1 coded as 0 1 0  
Class 2 coded as 0 0 1  
**NB:** the classes are **not coded** in a **binary code**.

The number of classes = 3

The corresponding formula of **SoftMax** function is :

$$f(x_i) = \frac{e_i^x}{\sum_{j=0}^N e_j^x}$$

$N$  is the Number of samples <sup>15</sup>





## 4- Single Layer Perceptron examples

### Example with neurolab: delta rule

```
# separate the data into test and train samples
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X,y,test_size=0.2)
```

```
# we have to construct the target values (true labels)
# we will convert each of the labels vectors into a 3 columns
# vector
from sklearn.preprocessing import LabelBinarizer
from sklearn.utils import shuffle
myLB = LabelBinarizer()
y3_train= myLB.fit_transform(y_train)
y3_test= myLB.fit_transform(y_test)
```

```
# train the SLP
myErr2 = mySLPDelta.train(x_train,y3_train, epochs=1000, show=100, lr=0.01)
```

```
# define a prediction function
def Predict(x, Net):
    if np.ndim(x)==1:
        x=[x]
    res = Net.sim(x)
    return np.argmax(res,axis=1)
```

Select the index corresponding to the **maximum probability** as the **predicted class label**

[By Amina Delali]

```
# compute the accuracy of the SLP predictions
# for the training and testing data
```

```
yPred = Predict(x_train,mySLPDelta)
accuracy =np.count_nonzero(yPred== y_train) / y_train.shape[0]
print ("Accuracy of the training = ",np.round(accuracy,2))
```

```
yP_test = Predict(x_test,mySLPDelta)
accuracy2 =np.count_nonzero(yP_test== y_test) / y_test.shape[0]
print ("Accuracy of the testing = ",np.round(accuracy2,2))
```

Accuracy of the training = 0.97  
Accuracy of the testing = 0.97

array([2, 0])

array([[0, 0, 1],  
[0, 0, 1],  
[1, 0, 0]])

**Third class**  
→ **third column == 1**

This implementation in **neurolab** uses the **sum squared error** as cost function. And it **doesn't** use the **derivative** of the activation function used

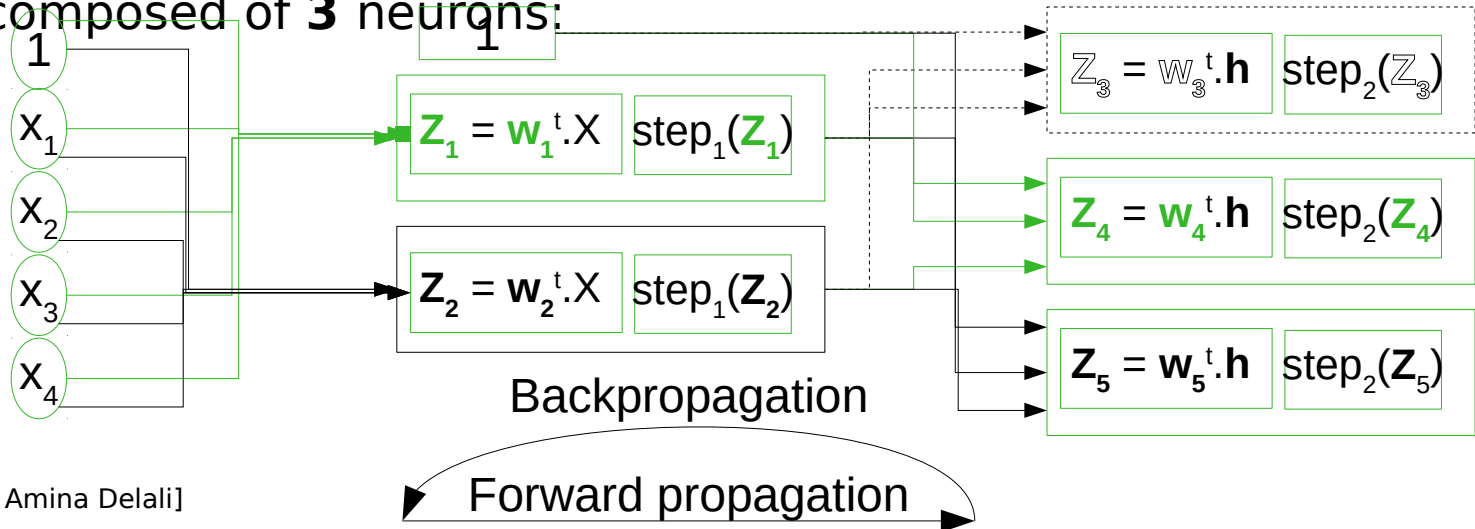




## 5- Multi Layer Perceptron

### Multi-Layer Perceptron

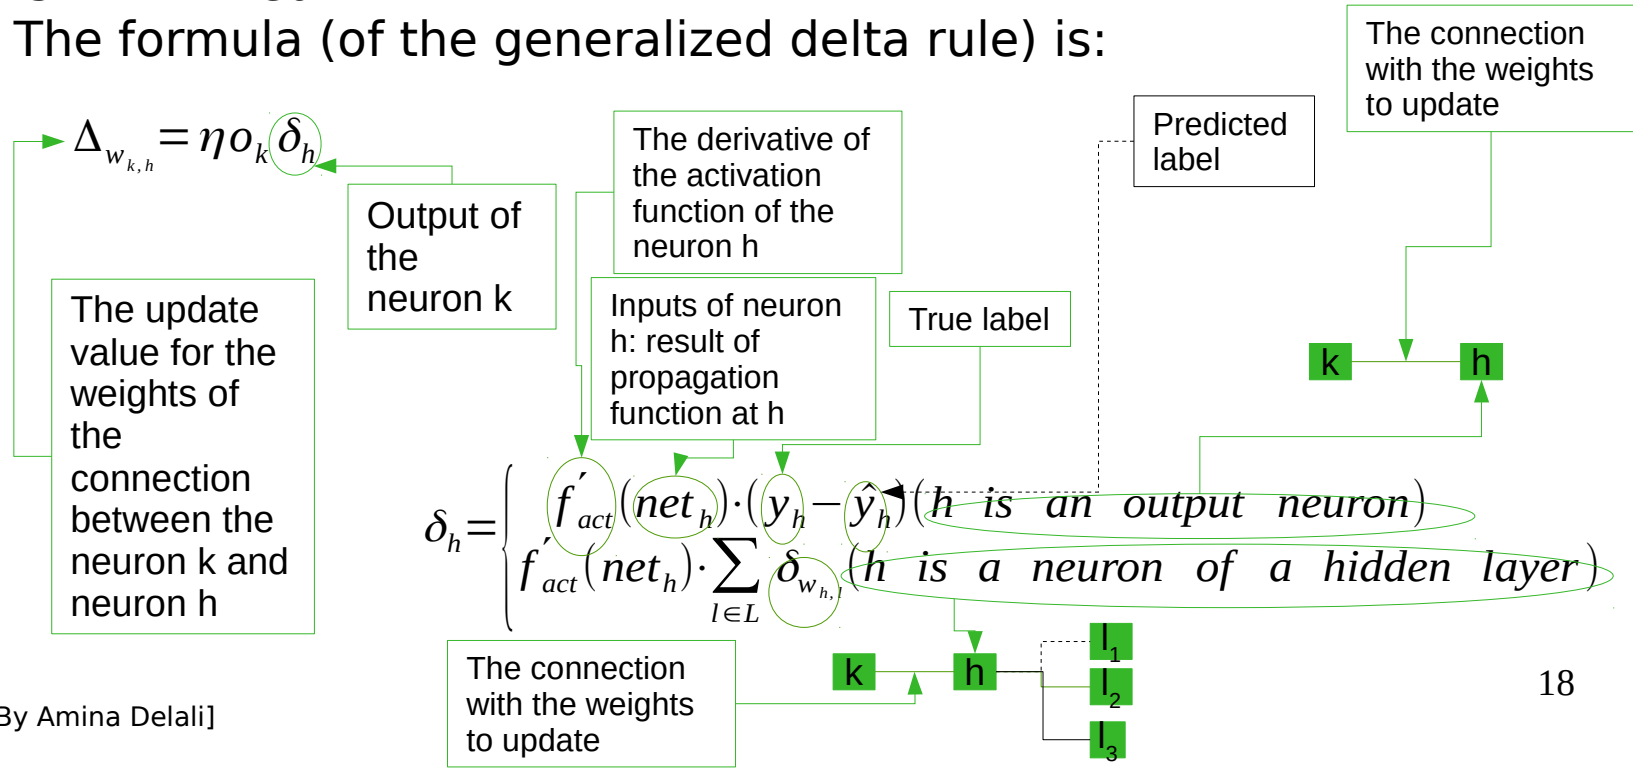
- An **MLP** (**M**ulti-**L**ayer **P**erceptron) is a **Percetron** with **one** or **more hidden layers**.
- It is another **Feed Forwad Artificial neural network**. Each of the layers (**except** the **output** layer) includes a **bias** neuron.
- An **ANN** with more than **one hidden** layer is a **Deep Neural Network ( DNN )**.
- Lets consider this MLP: input layer with **4** neurons, second(**hidden**) layer with **2** neurons (for both first and second layer the bias neurons are not counted). The last layer is composed of **3** neurons:





## Backpropagation

- It is a generalization of the **delta rule**. After a **forward pass**, a **backward** pass is applied to **update** the **weights** to back propagate the **errors**, using **gradient descent** procedure.
- This forward/backward passes are repeated until the error function is minimized
- The formula (of the generalized delta rule) is:





## 5- Multi Layer Perceptron

### Example

```
1 # an MLP with an input layer with 4 neurons, a hidden layer with
2 # 2 neurons and an output layer with 3 neurons
3 myMLP = nl.net.newff([[f1_mi,f1_ma],[f2_mi,f2_ma],[f3_mi,f3_ma],[f4_mi,f4_ma]],[2, 3],
4                    [nl.trans.TanSig()], [nl.trans.SoftMax()])
```

Activation function  
for the second  
(first hidden) layer

Activation function  
for the third  
(output) layer

```
1 # train the MLP
2 myErr2 = myMLP.train(x_train,y3_train, epochs=10000, show=100)
```

The goal of learning is reached

TanSig formula:

$$y = \frac{2}{1 + e^{-2n}} - 1$$

```
1 # compute the accuracy of the MLP predictions
2 # for the training and testing data
3
4 yPred = Predict(x_train,myMLP)
5 accuracy = np.count_nonzero(yPred == y_train) / y_train.shape[0]
6 print ("Accuracy of the training = ",np.round(accuracy,2))
7
8 yP_test = Predict(x_test,myMLP)
9 accuracy2 = np.count_nonzero(yP_test == y_test) / y_test.shape[0]
10 print ("Accuracy of the testing = ",np.round(accuracy2,2))
```

Accuracy of the training = 1.0  
Accuracy of the testing = 0.93

The  
accuracy of  
training is  
equal = 1.0  
(100%) but  
with test data  
we got less  
good results  
as with the  
SLP ==>  
over-fitting




## 6- ANN topologies

### FeedForward Neural Networks


- We have already seen **feedforward neural networks** (SLP and MLP)
  - One input layer + n hidden layers + one output layer ( $n \geq 1$ )
  - Connections are only allowed to neurons of the following layers
  - They can have **shortcut connections**: the connections are not set to the following layers but to subsequent layers

The white squares will correspond to the existing connections



Connections from the First neuron: it has only connections to the third and fourth neuron


The corresponding Hinton Diagram






## Recurrent Networks

### Direct Recurrence





- Multiple layers with connections allowed to neurons of the following layers
- A neuron can also be connected to itself

For visualization purposes the recurrent connections are represented by smaller squares



### Indirect Recurrence




- Multiple layers with connections allowed to neurons of the following layers
- Connections are also allowed between neurons and preceding layer

## Fully connected Neural Network



### 6- ANN topologies

- Multiple layers with connections allowed from any neuron to any other neuron
- Direct recurrence is not allowed
- Connections must be symmetric



```
1 myFC = 49 * [3]
2 myFC = np.array(myFC).reshape(7,7)
3 for i in range(7):
4     myFC[i,i] = 0.01
5 #print(myFC)
6 plots.hinton(myFC, add_legend=False)
```



# References

- Joshi Prateek. Artificial intelligence with Python. Packt Publishing, 2017.
- Jake VanderPlas. Python data science handbook: essential tools for working with data. O'Reilly Media, Inc, 2017.
- Neural Networks – II, Version 2 CSE IIT, Kharagpur, available at: <https://nptel.ac.in/courses/106105078/pdf/Lesson%2038.pdf>
- Isaac Changhau, Loss Functions in Neural Networks, 2017, available at: [https://isaacchanghau.github.io/post/loss\\_functions/](https://isaacchanghau.github.io/post/loss_functions/)
- Sebastian Seung, The delta rule, MIT Department of Brain and Cognitive Sciences , Introduction to Neural Networks, 2005, [https://ocw.mit.edu/courses/brain-and-cognitive-sciences/9-641j-introduction-to-neural-networks-spring-2005/lecture-notes/lec19\\_delta.pdf](https://ocw.mit.edu/courses/brain-and-cognitive-sciences/9-641j-introduction-to-neural-networks-spring-2005/lecture-notes/lec19_delta.pdf)
- NeuPy, Neural Networks in Python, <http://neupy.com/pages/home.html>



# Thank you!

FOR ALL YOUR TIME