



Ensemble Learning: Random Forests

AAA-Python Edition



school of a

Plan

- 1- Random Forest
- 2- Extreme Random Forest
- 3- Balanced Random Forest
- 4- Feature Importance
- 5- Grid Search
- 6- Balancing by resampling

Concept

- A random forest is :

Bagging or pasting.

+

Decision Trees:
with random
split

==

Bagging with **Decision Trees** as classifier. The trees use a random subset of features at each split

==

An **ensemble learning method** composed only of Decision Trees. They are trained on different random subsets of training data (with or without replacement) . To split a node, the trees select the best features from a random subset of features .



school of ants

1- Random Forest



school of ai

1- Random Forest

Random forest in sklearn: method 1

- Using a Bagging Classifier: since a random forest is a special case, of a bagging method. We can use a Bagging Classifier with decision trees as estimator.

```
1 from sklearn.ensemble import BaggingClassifier
2 from sklearn.tree import DecisionTreeClassifier
3 myEstimator= DecisionTreeClassifier(splitter="random",max_features="auto",max_depth=4)
4 myRF1 = BaggingClassifier(myEstimator,n_estimators=500,
5                           max_samples=1.0, bootstrap=True, n_jobs=-1,oob_score=True)
6
7 myRF1.fit(x_train,y_train)
8 y_pred= myRF1.predict(x_test)
9 print("Confusion matrix = \n",confusion_matrix(y_test,y_pred))
10 print("score=", np.round(myRF1.score(x_test,y_test),2))
```

Confusion matrix =
[[11 0 0]
[0 10 0]
[0 1 16]]
score= 0.97

By specifying the splitter= "random", at each split, the chosen feature will be the **"best random feature"**: the best feature from a random subset of features.



school of ai

1- Random Forest

Random forest in sklearn: method 2

- Using the random forest classifier class: in sklearn, the size of the samples used in the trees is the same as the original data size.

```
1 from sklearn.ensemble import RandomForestClassifier
2
3 myRF2= RandomForestClassifier(n_estimators=500, n_jobs=-1, max_depth=4, oob_score=True)
4 myRF2.fit(x_train, y_train)
5 y_pred = myRF2.predict(x_test)
6
7 print("Confusion matrix = \n", confusion_matrix(y_test, y_pred))
8 print("score=", np.round(myRF2.score(x_test, y_test), 2))
```

Confusion matrix =
[[11 0 0]
[0 10 0]
[0 1 16]]
score= 0.97

Same parameters
used in the bagging
classifier for the
previous example.

There is no “splitter” parameter. Almost all other default parameters are the same. The only one that differs is : max_features which is here set by default to “auto”. The only parameter we had to specify is (to be adequate with the tree of the previous example) is : max_depth

max_features = “auto” means that:
max_features = sqrt(num_features)

Concept



school of a

2- Extreme Random Forest

- In **Extreme random forests**, the **thresholds** used in the decision trees are **selected randomly**, instead of choosing the best one.
- They are also called “**Extremely randomized Trees**” or “**Extra Trees**”.
- In sklearn, they are: “**ExtraTreesClassifier**” and “**ExtraTreesRegressor**”
- “Extra**Trees**Classifier” are random forests, but there is in sklearn “Extra**Tree**Classifier” which is a decision tree and not a forest.



school of a

2- Extreme Random Forest

Example

```
1 from sklearn.tree import ExtraTreeClassifier
2
3 myEstimator2= ExtraTreeClassifier(splitter="random",max_features="auto",max_depth=4)
4 myERF1 = BaggingClassifier(myEstimator2,n_estimators=500,
5                             max_samples=1.0, bootstrap=True, n_jobs=-1,oob_score=True)
6
7 myERF1.fit(x_train,y_train)
8 y_pred = myERF1.predict(x_test)
9
10 print("Confusion matrix = \n",confusion_matrix(y_test,y_pred))
11 print("score=", np.round(myERF1.score(x_test,y_test),2))
```

Confusion matrix =
[[11 0 0]
[0 10 0]
[0 1 16]]
score= 0.97

```
1 from sklearn.ensemble import ExtraTreesClassifier
2
3 myERF2= ExtraTreesClassifier(n_estimators=500, n_jobs=-1,max_depth=4,oob_score=True,bootstrap=True)
4 myERF2.fit(x_train, y_train)
5 y_pred = myERF2.predict(x_test)
6
7 print("Confusion matrix = \n",confusion_matrix(y_test,y_pred))
8 print("score=", np.round(myERF2.score(x_test,y_test),2))
```



school of ants

2- Extreme Random Forest

Out of bag scores

- The bag refers to the training samples selected and used by a predictor in a bagging method.
- The out of bag, are the remaining samples.
- For each predictor, it is possible to compute its score using its out of bag samples.
- At the end of the training, the average of all these scores will represent the “out of bag score” of the learning method.
- In sklearn, you have to specify: “oob_score = True” as parameter of the bagging method.

```
myRF1 = BaggingClassifier(myEstimator, n_estimators=500,  
                          max_samples=1.0, bootstrap=True, n_jobs=-1, oob_score=True)
```

```
myRF2 = RandomForestClassifier(n_estimators=500, n_jobs=-1, max_depth=4, oob_score=True)
```

```
print("Out of bag score:", np.round(myRF1.oob_score_, 2))
```

Out of bag score: 0.94

```
print("Out of bag score:", np.round(myRF2.oob_score_, 2))
```

Out of bag score: 0.95

```
print("Out of bag score:", np.round(myERF1.oob_score_, 2))
```

```
print("Out of bag score:", np.round(myERF2.oob_score_, 2))
```

Out of bag score: 0.96

Concept

- Sometimes, in classification, the representation of your classes in the training data is “**unbalanced**”. For example, in a binary classification, you can have a big number of samples belonging to the first class. And, only few samples related to the other one.
- In this case, the learning method will not be able to classify correctly.
- To avoid this situation, scikit-learn allow us to define a parameter called: **class_weight**, with a special value: “**balanced**” ==> it will attribute bigger weights for less present classes.
- The formula is as follow: $n_samples / (n_classes * np.bincount(y))$

Count number of occurrences of each value in y



school of ai

3- Balanced Random Forest



school of ai

3- Balanced Random Forest

Example without balancing

```
from pandas import DataFrame as DF, Series as S
import pandas as pd
myData = pd.read_csv("AAA-Ped-Week5/A3P-w5-data_imbalance.txt", header = -1)

x2= myData.iloc[:,[0,1]].values
y2= myData.iloc[:,2].values
```

| | 0 | 1 | 2 |
|---|------|------|---|
| 0 | 5.66 | 6.77 | 1 |
| 1 | 4.40 | 5.05 | 0 |
| 2 | 3.52 | 4.73 | 1 |

Loading the
data, and
extracting x
and y values

We will use **yellowbrick** library to visualize
our classification boundaries.

```
from yellowbrick.contrib.classifier import DecisionViz


myRF2= RandomForestClassifier(n_estimators=100, n_jobs=-1, oob_score=True, max_depth=4,
                             )
myUVis = DecisionViz(myRF2, "Unbalanced Learning", features = ["Feature 1", "Feature 2"],
                     classes = [0,1])

myUVis.fit(x2_train,y2_train)
myUVis.draw(x2_test, y2_test)
myUVis.poof(outpath="unbalanced.png")
```



3- Balanced Random Forest

Example without balancing (suite)



The classes are not balanced

We see that some samples belonging to class 0 are predicted to be in class 1 and vice versa.

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.65 | 0.66 | 0.66 | 62 |
| 1 | 0.93 | 0.93 | 0.93 | 313 |



school of ai

3- Balanced Random Forest

Same example with balancing


```
myBRF2= RandomForestClassifier(n_estimators=100, n_jobs=-1, oob_score=True, max_depth=4,  
                               class_weight="balanced")  
myVis = DecisionViz(myBRF2, "Balanced Learning", features = ["Feature 1", "Feature 2"],  
                    classes = [0,1])  
  
myVis.fit(x2_train,y2_train)  
myVis.draw(x2_test, y2_test)  
myVis.poof(outpath="balanced.png")
```

Balanced classes

precision recall f1-score support

| | | | | |
|---|------|------|------|-----|
| 0 | 0.53 | 0.92 | 0.67 | 62 |
| 1 | 0.98 | 0.84 | 0.91 | 313 |

- We see clearly that the decision boundaries changed: the area of class 0 is bigger now
- But, we have more classes 1 predicted as belonging to class 0.



[By Amina Delali]

Concept

- **GridSearch** is a tool used to test a learning model with a list of dictionaries containing different parameters.
- We will use it to find the best parameters for a **Random Forest** method. We will apply it on the already used **iris** data.

```
# variate n_estimators, max_depth; and min_samples_leaf
myGrid = [ {'n_estimators': [20,40,60], "min_samples_leaf": [1,2,3]},
           { 'n_estimators': [10,20,30,40], "max_depth": [1,2,3]}
]

from sklearn.model_selection import GridSearchCV
# A grid search instantiation using a Random Forest classifier
myEstimator = RandomForestClassifier()
myClassifier = GridSearchCV(estimator= myEstimator,
                           param_grid= myGrid)
myClassifier.fit(x_train, y_train)
```

```
1 print("\nThe obtained scores are:")
2 for params, mean, allv in myClassifier.grid_scores_:
3     print(params, '-->', np.round(mean, 3))
4 print("The best parameters are :", myClassifier.best_params_)
```

First the first dictionary we have: 9 combinations of parameters. For the second one we have: 12 combinations.



school of a

4- Grid Search

Example results

```
{ 'min_samples_leaf': 1, 'n_estimators': 20 } ==> 0.964
{ 'min_samples_leaf': 1, 'n_estimators': 40 } ==> 0.964
{ 'min_samples_leaf': 1, 'n_estimators': 60 } ==> 0.964
{ 'min_samples_leaf': 2, 'n_estimators': 20 } ==> 0.973
{ 'min_samples_leaf': 2, 'n_estimators': 40 } ==> 0.973
{ 'min_samples_leaf': 2, 'n_estimators': 60 } ==> 0.964
{ 'min_samples_leaf': 3, 'n_estimators': 20 } ==> 0.973
{ 'min_samples_leaf': 3, 'n_estimators': 40 } ==> 0.973
{ 'min_samples_leaf': 3, 'n_estimators': 60 } ==> 0.964
{ 'max_depth': 1, 'n_estimators': 10 } ==> 0.946
{ 'max_depth': 1, 'n_estimators': 20 } ==> 0.955
{ 'max_depth': 1, 'n_estimators': 30 } ==> 0.964
{ 'max_depth': 1, 'n_estimators': 40 } ==> 0.955
{ 'max_depth': 2, 'n_estimators': 10 } ==> 0.955
{ 'max_depth': 2, 'n_estimators': 20 } ==> 0.964
{ 'max_depth': 2, 'n_estimators': 30 } ==> 0.964
{ 'max_depth': 2, 'n_estimators': 40 } ==> 0.973
{ 'max_depth': 3, 'n_estimators': 10 } ==> 0.964
{ 'max_depth': 3, 'n_estimators': 20 } ==> 0.964
{ 'max_depth': 3, 'n_estimators': 30 } ==> 0.964
{ 'max_depth': 3, 'n_estimators': 40 } ==> 0.964
```

- We can see that augmenting the number of estimators doesn't necessarily enhance the results.
- Same observation about max_depth value

The best parameters are : { 'min_samples_leaf': 2, 'n_estimators': 20 }



school of ai


4- Grid Search

Best result visualization

- We will display the confusion matrix using **yellowbrick**.

```
1 from yellowbrick.classifier import ConfusionMatrix
2 myEstimator= RandomForestClassifier(min_samples_leaf=2,n_estimators= 20)
3 myVisualizer = ClassificationReport(myEstimator, classes=["setosa","versicolour","virginica"],
4                                     support=True)
5 myVisualizer.fit(x_train,y_train)
6 myVisualizer.score(x_test,y_test)
7 myVisualizer.poof(outpath="Confusion_
```

The confusion matrix is printed as a "heat-map": the best values are the darkest, and vice-versa.



Concept

- The **features** in a data set have **not** the **same importance**.
- In a **decision tree**, this importance can be deduced from the **distance** between the **appearance** of a **feature** in a **decision node** and **the root**.
- In a **random forest**, the **average** of the distances corresponding to each tree will represent the **feature's importance**.
- **Scikit-learn** implements this method.

```
1 # we will use the previous trained estimator
2
3 for feature, importance in zip(myIris["feature_names"], myEstimator.feature_importances_):
4     print(feature, '==>', np.round(importance,5))
5
```

sepal length (cm) ==> 0.07181
sepal width (cm) ==> 0.00342
petal length (cm) ==> 0.3492
petal width (cm) ==> 0.57557

We can see clearly that the most important features are: "petal length" and "petal width"






5- Feature Importance

Visualization of the features importance

- Again, we will use **yellowbrick**.

```
1 from yellowbrick.features import FeatureImportances
2
3 #features names
4 features = myIris["feature_names"]
5 # instantiate the feature importance visualizer
6 FIVisualizer = FeatureImportances(myEstimator, labels=features)
7 # train the visualizer
8 FIVisualizer.fit(X, y)
9 # display the relative features importance
10 FIVisualizer.pooof("feature_importance.png")
11
```

Most important
features






school of a

5- Feature Importance

Comparison of results


```
myVisualizer.fit(x_train[:,[2,3]].reshape(-1,2),y_train)
```



Using only **2 most important** features
(same result as using
all the features)

[By Amina Delali]

```
myVisualizer.fit(x_train[:,[0,1]].reshape(-1,2),y_train)
```



Using only **2 least important** features
(same result as using
all the features)

Concept



school of AI

6- Balancing by resampling

- An other way to obtain a balanced training set, is to **remove** samples from the majority class. Which is called: **under-sampling**. Or, to **add** new samples belonging to the less represented class. Which is called: **over-sampling**.
- There is a variety of resampling techniques. The library **imbalanced-learn** library implements some of them.
- For a **RandomForest classifier**, it applies: random under-sampling technique with different strategies.
- The default one is to : resample all classes but the minority class;



school of ai

6- Balancing by resampling

Example

- In this example, we combine the **2** libraries: **yellowbrick** and **imbalanced-learn**


```
1 from imblearn.ensemble import BalancedRandomForestClassifier
2 from sklearn.metrics import balanced_accuracy_score
3
4 myBRF3 = BalancedRandomForestClassifier(n_estimators=100, n_jobs=-1, oob_score=True, max_depth=4)
5 myVis3 = DecisionViz(myBRF3, "Balanced Learning2", features = ["Feature 1", "Feature 2"],
6                       classes = [0,1])
7
8
9 myVis3.fit(x2_train, y2_train)
10 myVis3.draw(x2_test, y2_test)
11 myVis3.poof(outpath="balanced2.png")
12
13 print(classification_report(y2_test, myBRF3.predict(x2_test)))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.50 | 0.89 | 0.64 | 72 |
| 1 | 0.97 | 0.79 | 0.87 | 303 |
| micro avg | 0.81 | 0.81 | 0.81 | 375 |
| macro avg | 0.73 | 0.84 | 0.75 | 375 |
| weighted avg | 0.88 | 0.81 | 0.82 | 375 |

Example (suite)

We can see that the **class 0** region obtained is different than the one obtained using scikit-learn weighting technique

```
3 from IPython.display import Image
4 Image(filename='balanced2.png')
```





school of ai

References

- Aurélien Géron. Hands-on machine learning with Scikit-Learn and Tensor-Flow: concepts, tools, and techniques to build intelligent systems. O'Reilly Media, Inc, 2017.
- imbalanced learn. User guide. On-line at https://imbalanced-learn.org/en/stable/user_guide.html. Accessed on 09-12-2018.
- Joshi Prateek. Artificial intelligence with Python. Packt Publishing, 2017.
- Alencar Rafael. Resampling strategies for imbalanced datasets. On-line at <https://www.kaggle.com/rafjaa/resampling-strategies-for-imbalanced-datasets>. Accessed on 09-12-2018.
- Scikit-learn.org. scikit-learn, machine learning in python. On-line at <https://scikit-learn.org/stable/>. Accessed on 03-11-2018.
- yellowbrick. Yellowbrick: Machine learning visualization. On-line at <http://www.scikit-yb.org/en/latest/>. Accessed on 09-12-2018.



Thank you!

FOR ALL YOUR TIME